

Vmem- φ : Zero-Cost Out-of-Distribution Detection in Spiking Neural Networks from Membrane-Potential Statistics

Anonymous authors
Paper under double-blind review

Abstract

Spiking neural networks (SNNs) on event cameras are deployed precisely where the standard out-of-distribution (OOD) toolkit cannot run: on neuromorphic silicon with no floating-point logits, no gradients, and no budget for a second forward pass. We observe that the sub-threshold membrane potential $V_{\text{mem}}(t)$ of the Parametric Leaky-Integrate-and-Fire (PLIF) neurons is a free byproduct of the forward pass – a register the substrate already maintains – and that its per-channel statistics $\varphi = [\mu, \sigma^2, \kappa]$ (a 2112-D descriptor over 704 channels $\times$ 4 layers) form an OOD signal at zero additional compute. We contribute Gen1-C, a reproducible, seeded event-camera corruption benchmark of six sensor-grounded corruptions $\times$ five severities – each traced to a concrete hardware failure mechanism – plus a released membrane-statistics feature set so OOD detectors can be compared without re-running SNN inference. We give a closed-form leaky-integrator theory that predicts, per corruption, which membrane moment moves and therefore which corruptions are detectable and which are provably not, and we confirm with a motivation experiment that feeding these corruptions to a pretrained state-of-the-art hybrid SNN-ANN detector (no retraining) collapses its mAP. Finally we introduce the Manifold-Decomposition Detector (MDD), which decomposes each frame’s deviation into orthogonal manifold axes scored two-sided, taking the contraction corruptions above chance for the first time; one unsupervised, corruption-blind score reaches ≥ 0.96 AU-ROC on four of six corruptions with sequence aggregation. We report honest negatives – polarity_flip (a network invariance) and spatial_dropout (a signature global pooling discards) – and frame them as an open challenge: the Achilles’ heel of event-based perception.

1 Introduction

Event cameras report per-pixel brightness changes asynchronously at microsecond resolution, high dynamic range, and low power; paired with spiking neural networks (SNNs) they form the canonical low-power perception stack for neuromorphic hardware. As these systems move toward safety-critical deployment (automotive, robotics), a basic question becomes unavoidable: when the sensor input drifts away from the training distribution, can the system notice? Out-of-distribution (OOD) detection is the standard answer for artificial neural networks (ANNs), but essentially every method in that literature assumes a capability the neuromorphic execution model does not provide – floating-point logits (MSP, Energy), input gradients (ODIN, GradNorm), or a second forward pass over the features (ReAct, ANN-feature Mahalanobis/kNN). On a chip that emits spikes and integer events under a strict energy budget, these are not merely expensive but architecturally infeasible (Section 3.4).

We start from a signal that is free exactly where the alternatives are impossible. The Parametric Leaky-Integrate-and-Fire (PLIF) neuron maintains a sub-threshold membrane potential $V_{\text{mem}}(t)$ – physical state the substrate already updates on every time step to decide whether to spike. Reading it and reducing it to a few per-channel moments adds

zero multiply-accumulates and zero weight fetches. We call the resulting 2112-dimensional descriptor φ and ask how far this free signal goes as an OOD detector against realistic event-camera corruptions, each grounded in a concrete sensor failure mechanism (Figure 1).

Contributions. We make three key contributions.

1. A zero-cost membrane OOD signal and the Gen1-C benchmark. We introduce φ , an OOD signal harvested from the PLIF membrane at zero additional compute, and argue it is the only OOD method that runs inside the compute envelope of the target neuromorphic hardware (Section 3.4). To evaluate it we release Gen1-C: a reproducible (seeded, deterministic, model-agnostic) event-camera corruption benchmark of six corruptions $\times$ five monotonic severities, each mapped to a physical hardware failure mode (Section 3.2, Figure 1), in the ImageNet-C (Hendrycks & Dietterich, 2019) tradition, plus a released membrane-statistics feature set.
2. A closed-form theory of detectability, empirically motivated. We derive, from the PLIF state equation and with no parameters fit to corruption labels, a leaky-integrator theory that predicts which membrane moment each corruption moves and therefore which corruptions are detectable – a dilation/contraction/invariant taxonomy (Section 3.3). We confirm the stakes with a motivation experiment: a pretrained SOTA hybrid SNN-ANN detector (Neftci et al., 2025) loses substantial mAP under these corruptions with no retraining (Section 4.2).
3. The Manifold-Decomposition Detector (MDD), and an open challenge. Following the geometric diagnosis, MDD stops collapsing φ to one scalar and instead scores orthogonal, two-sided manifold axes fused by a calibrated OR (Section 3.5); it takes the contraction corruptions above chance for the first time and reaches ≥ 0.96 AUROC on four of six corruptions with sequence aggregation. We report `polarity_flip` and `spatial_dropout` as honest, theory-predicted residuals and release them as a community challenge (Section 5).

2 Background and Related Work

OOD detection for ANNs. The dominant families score a post-hoc statistic of a trained network: maximum softmax probability (Hendrycks & Gimpel, 2017), temperature-scaled and input-perturbed logits (Liang et al., 2018), free energy (Liu et al., 2020), rectified activations (Sun et al., 2021), virtual-logit matching (Wang et al., 2022), gradient norms (Huang et al., 2021), and distance/density on penultimate features (Mahalanobis (Lee et al., 2018), deep k -NN (Sun et al., 2022)). All presuppose floating-point readouts, gradients, or repeated forward passes – capabilities absent from the spiking, feed-forward, integer-event execution model of neuromorphic hardware (Section 3.4 makes this precise).

OOD detection for SNNs (and how we differ). OOD detection inside an SNN is not unprecedented. Most directly, Martínez et al. (Martínez et al., 2023) characterise the internal activations of hidden layers via spike-count patterns and add an explainability map, evaluated on image (including event-based) classification benchmarks. Our work is distinct on four axes that, in combination, define a new task: (i) we read the sub-threshold membrane potential, not spike counts – a continuous, pre-threshold signal with strictly more information than the binarised spikes; (ii) the host model is an event-camera object detector (Prophesee Gen1), not a classifier, so there is no class posterior to lean on; (iii) the distribution shifts are event-stream corruptions we define and release as a benchmark, not held-out classes; and (iv) we make the zero-compute / neuromorphic-register feasibility argument central. We therefore position the paper as introducing the event-camera membrane-statistics OOD task and the Gen1-C benchmark, while explicitly crediting prior SNN-OOD work.

Event-camera robustness and corruption benchmarks. Robustness benchmarks built from synthetic corruptions are standard for frame cameras (ImageNet-C (Hendrycks & Dietterich, 2019)); the event-camera analogue is comparatively underdeveloped. Gen1-C (Section 3.2) fills this gap with corruptions grounded in concrete sensor failure modes (Figure 1, Table 1),

seeded for exact reproducibility, and operating on the standard histogram tensor so they compose with any event model.

The host detector. Our membrane signal is read from the pretrained hybrid SNN-ANN detector of Neftci et al. (Neftci et al., 2025), a strong Gen1 baseline that fuses a spiking front-end (SpikingJelly PLIF (Fang et al., 2021)) with an ANN detection head. We treat it as read-only and never retrain it – both for the OOD signal and for the motivation experiment of Section 4.2.

Membrane potential as a hardware register. On Loihi (Davies et al., 2018), BrainScaleS (Schemmel et al., 2010), and SpiNNaker (Furber et al., 2014), V_{mem} is state the neuron already maintains; harvesting φ from it is an $O(\text{channels})$ read. This is the basis of the feasibility argument in Section 3.4.

3 Method

We present (i) the zero-cost membrane descriptor φ , (ii) the Gen1-C corruption suite with its physical-mechanism grounding and algorithms, (iii) a leaky-integrator theory that predicts per-corruption detectability, (iv) the neuromorphic-hardware feasibility argument, and (v) the Manifold-Decomposition Detector that follows from the theory.

3.1 The zero-cost membrane descriptor φ

The membrane signal is read from the pretrained hybrid SNN-ANN detector with a forward hook on all four PLIF layers (704 channels total). For each channel c we apply Global Average Pooling (GAP) over the spatial dimensions of the sub-threshold trace V and compute its first three central moments across the T SNN time steps,

$$\varphi_c = [\mu_c, \sigma_c^2, \kappa_c], \quad \mu_c = \mathbb{E}[V_c], \quad \sigma_c^2 = \text{Var}[V_c], \quad \kappa_c = \frac{\mathbb{E}[(V_c - \mu_c)^4]}{\sigma_c^4}, \quad (1)$$

stacking to a 2112-D descriptor φ per frame. The mean captures sustained drive, the variance fluctuation energy, the kurtosis burstiness. Two properties drive every later result: (i) GAP discards spatial layout – φ knows how much each channel fired, not where; (ii) φ inherits the network’s invariances. A spatial-dispersion extension phi_spatial (1408-D: per-channel spatial variance and participation ratio) recovers part of the structure GAP discards. Clean φ lies on a thin, curved, scene-dependent manifold; detection is deciding whether a new φ lies on it.

3.2 The Gen1-C event-camera corruption suite

A benchmark for OOD detection is only as meaningful as the distribution shifts it tests against. We define a controlled, reproducible battery of six corruptions that model the dominant failure modes of event-camera front-ends, each at five monotonically increasing severities; with the clean baseline this yields $6 \times 5 + 1 = 31$ runs. Each corruption perturbs a structurally distinct axis of the input – DC offset, localized burst, temporal order, global rate, polarity, or spatial support – which is what makes the closed-form analysis of Section 3.3 possible.

Input representation. Every corruption operates on the model’s input tensor, so it composes with any downstream architecture and is label-free. A sequence is a stack of N event-histogram frames of shape $(20, H, W)$, $H \times W = 240 \times 304$ for Gen1. The 20 channels are two polarities $\times$ ten $50 \mu\text{s}$ time bins: channels 0–9 ON (oldest to newest), 10–19 OFF. Counts are uint8; corruptions promote to a wider type internally and clip to $[0, 255]$.

Design principles. (i) Reproducibility: every random decision is drawn from a single seeded NumPy generator, so a (corruption, severity, seed) triple is deterministic. (ii) Hardware agnosticism: operations dispatch to NumPy or CuPy. (iii) Severity monotonicity: severity 1 is mild and 5 severe, with monotone parameters (Table 2), enabling the Spearman- ρ analysis of Section 4.



Figure 1: The six Gen1-C corruptions and their physical mechanisms. For each corruption: a normal event frame (left) and the corrupted frame the modelled hardware fault produces (right). `hot_pixel` adds persistent stuck-pixel firings; `event_flood` injects a dense local burst; `temporal_jitter` smears events across time bins; `event_rate_shift` scales the global density; `polarity_flip` swaps ON/OFF; `spatial_dropout` blanks a region. Each maps to a distinct, predictable signature in the membrane moments (Section 3.3).

3.2.1 Physical mechanisms: how each corruption arises in hardware

The six corruptions are not arbitrary perturbations; each models a documented event-camera failure, which is what makes Gen1-C a robustness test rather than a synthetic curiosity. Figure 1 contrasts the normal sensor output with the corrupted output for each, and names the mechanism; we expand the physics here.

- `hot_pixel` – defective / stuck pixels, fixed-pattern noise. Fabrication defects, bias-current mismatch and junction leakage make individual DVS pixels cross their change threshold with no scene change, firing continuously at fixed coordinates. The rate of such self-triggering rises with temperature and with aggressive bias settings, so it materializes as a constant DC offset at a fixed set of locations – the event-camera analogue of stuck CCD/CMOS pixels.

Corruption	Physical failure mode	Axis of the input perturbed
hot_pixel	defective / stuck pixels	constant DC offset at fixed pixel locations
event_flood	EMI / sensor saturation	localized count bursts in space and time
temporal_jitter	clock noise / timestamp error	time-bin ordering (per polarity)
event_rate_shift	event-density domain gap	global multiplicative rate
polarity_flip	ON/OFF misclassification	polarity sign on a fraction of frames
spatial_dropout	dead-pixel clusters / occlusion	spatial support (rectangular regions zeroed)

Table 1: The six corruptions, the sensor failure each models, and the input axis it perturbs.

Corruption	Sev. 1	Sev. 2	Sev. 3	Sev. 4	Sev. 5
hot_pixel (n, δ)	10, 20	30, 40	80, 80	150, 140	300, 200
event_flood (b, f, δ)	2, .05, 60	4, .08, 100	8, .12, 150	15, .18, 200	25, .25, 255
temporal_jitter $s_{\max}$	1	2	3	5	8
event_rate_shift (u, o)	.80, 1.20	.65, 1.40	.50, 1.65	.35, 2.00	.20, 2.50
polarity_flip p	.05	.10	.20	.35	.50
spatial_dropout (k, f_h, f_w)	1, .05, .05	1, .10, .10	2, .10, .10	2, .15, .15	3, .20, .20

Table 2: Severity schedules. Symbols: n hot pixels, δ injected count, b burst centres, f patch fraction, $s_{\max}$ max bin shift, (u, o) under/over rate factors, p flipped-frame fraction, k dropout regions, (f_h, f_w) region height/width fractions.

- event_flood – EMI, illumination transients, readout saturation. Electromagnetic interference, LED/PWM flicker, or a sudden strong-light transient drives a dense burst of spurious events in a region for a short window; when the local event rate exceeds the readout/arbitrator bandwidth the periphery saturates, so the burst is both localized in space and bounded in time.
- temporal_jitter – clock noise, timestamp error. Row/column arbitration delay, bus contention, and finite timestamp resolution misassign events to neighbouring time bins. The marginal per-pixel counts are essentially preserved; what is corrupted is the order (and hence the temporal autocorrelation) of events.
- event_rate_shift – contrast-threshold / biasing drift, domain gap. The per-pixel contrast threshold and bias currents set the event rate; temperature drift, re-biasing, or simply moving between a low-texture/slow-motion and a high-texture/fast-motion regime multiplies the global rate. We model both directions (under- and over-rate) with a seeded coin flip.
- polarity_flip – ON/OFF misclassification. Comparator offset and asymmetric ON/OFF thresholds (a documented DVS non-ideality) cause the brightness-change sign to be reported incorrectly for a fraction of events; under fast global illumination swings whole frames can have their polarity inverted.
- spatial_dropout – dead-pixel clusters, occlusion. Manufacturing dead regions, dust/lens contamination, condensation, or partial physical occlusion zero the event support over contiguous areas for the duration of the fault.

3.2.2 Corruption algorithms and severity schedules

Each corruption operates on a histogram $X \in \{0, \dots, 255\}^{N \times 20 \times H \times W}$ and returns a corrupted copy. Hot pixel: add a constant δ to every bin/frame at n fixed random locations (int16, clipped). Event flood: at b random frame/patch centres add δ inside a patch covering fraction f over frames within ± 2 of the centre. Temporal jitter: roll the ON block (ch. 0–9) and OFF block (ch. 10–19) along the time-bin axis by independent integer shifts in $[-s_{\max}, s_{\max}]$, zero-filling vacated bins. Event-rate shift: multiply all counts by a seeded under/over factor and clip. Polarity flip: swap the ON/OFF blocks on a random fraction p of frames. Spatial dropout: zero k rectangular regions of size (f_h, f_w) across all channels/frames. Severity schedules are in Table 2.

3.3 Why membrane statistics detect corruption: a leaky-integrator theory

We show that the detectability of each corruption is a closed-form consequence of the PLIF state equation, not an empirical accident. The membrane is a leaky exponential accumulator of the input current; each corruption perturbs that current in a structurally distinct way, and the perturbation propagates to a predictable change in the moments of φ .

State equation and closed form. SpikingJelly’s MultiStepParametricLIFNode (Fang et al., 2021) updates, per channel, $V[t] = (1 - \frac{1}{\tau})V[t-1] + WX[t]$, emitting a spike with reset when V crosses threshold θ . Writing $\lambda \equiv 1 - 1/\tau \in (0, 1)$ and unrolling from $V[-1] = 0$,

$$V[t] = \sum_{k=0}^t \lambda^k W X[t-k], \quad (2)$$

so V is a causal, leaky linear filter of the event stream with memory horizon $\sim \tau$. Two consequences: V carries the temporal structure of recent input (so temporally-structured corruptions leave a signature even when the marginal input is unchanged), and because $X \rightarrow V$ is linear, the moments of V inherit closed-form dependence on the moments of X .

Steady-state moments. For wide-sense-stationary input with mean m , variance s^2 , autocovariance $\gamma(j)$,

$$\mathbb{E}[V] = \frac{Wm}{1-\lambda} = \tau Wm, \quad \text{Var}[V] = \frac{(Ws)^2}{1-\lambda^2} + 2W^2 \sum_{j \geq 1} \gamma(j) \frac{\lambda^j}{1-\lambda^2}. \quad (3)$$

The mean tracks the input rate amplified by τ ; the variance tracks the input variance plus an autocovariance term that depends on the time ordering, not the marginal. These are the levers each corruption pulls.

Per-corruption predictions. `hot_pixel` adds a constant δ , shifting $\mathbb{E}[V]$ by $\tau W\delta$ until V saturates at θ : affected μ_c separate from the clean range, giving inevitable detection (AUROC ≈ 1 , monotone). `temporal_jitter` is a measure-preserving permutation: it leaves m, s^2 (hence $\mathbb{E}[V]$ and the first variance term) unchanged but destroys $\gamma(j)$ and the lag-1 autocorrelation of $V(t)$ – visible only to temporal statistics and to deeper layers (larger effective τ), so the L4 slice is the strongest single view. `event_rate_shift` scales $m \rightarrow \alpha m$, displacing every channel’s mean along one global-activity direction – detectable by a 1-D projection, with a phase transition once α leaves the clean band. `spatial_dropout` lowers the spatial input variance, hence σ_c^2 , while μ_c stays intact; but GAP pools away the spatial dimension before the moments are taken, so the surviving signal makes the membrane look quieter and more typical – under one-sided distance scoring the detector inverts (AUROC < 0.5). `event_flood` raises m nearly uniformly, the same signature a genuinely busy clean scene produces, giving near-chance per frame; what distinguishes it (spatially unstructured activity; a temporally consistent bias) lives in the spatial dispersion GAP removes and in sequence-level consistency. `polarity_flip`: the trained detector is approximately polarity-symmetric, so WX – and thus φ – is near-invariant; the AUROC ceiling is ≈ 0.5 by construction.

Taxonomy. This yields a falsifiable taxonomy that organizes the whole results section: magnitude corruptions (`hot_pixel`, `event_rate_shift`, `jitter-via-variance`) move φ away from clean and are caught by distance detectors; structure corruptions (`spatial_dropout`, the spatial part of `event_flood`) leave the marginal moments intact and require the spatial-dispersion features GAP discards or a rule that flags “too typical”; invariant corruptions (`polarity_flip`) are near-orthogonal to the membrane and are an input-side problem. Every prediction follows from (2)–(3) with no parameters fit to corruption labels.

3.4 The neuromorphic hardware advantage

Existing OOD detectors are designed for ANNs on von Neumann hardware, where logits, gradients, and repeated forward passes are cheap. SNNs are deployed for the opposite reason – to run within the energy budget of neuromorphic silicon – and there the standard OOD

toolkit is not merely expensive but architecturally infeasible. On Loihi (Davies et al., 2018), BrainScaleS (Schemmel et al., 2010), and SpiNNaker (Furber et al., 2014) the membrane potential V is physical state the substrate already maintains; producing φ is an $O(\text{channels})$ register read plus three moments – zero MACs, zero weight fetches. Table 3 shows every competing family needs a capability the spiking, feed-forward, integer-event model does not provide. The benchmark of Section 4 should therefore be read not as “ φ narrowly beats ANN baselines on a GPU” but as “ φ provides OOD detection inside the compute envelope of the deployment hardware, where the alternatives cannot be evaluated at all.” We measure φ in simulation (SpikingJelly on GPU); the cost claim is architectural – an on-silicon energy measurement is future work.

3.5 The Manifold-Decomposition Detector (MDD)

The theory classifies corruptions by how they move φ relative to the clean manifold: dilations push outward (hot_pixel, event_rate_shift; jitter in the deep-layer subspace), contractions pull toward the mode (spatial_dropout, event_flood) and thereby invert one-sided distance/density detectors, and invariant corruptions (polarity_flip) leave no membrane signature. No single signed scalar can face outward, inward, and into a subspace at once – which is why every prior single detector (Mahalanobis, k -NN, GMM, flow, learned fusion) topped out near 0.70 macro and inverted on contractions. MDD stops collapsing φ to one number and scores three (optionally four) orthogonal, two-sided axes, fused by a calibrated OR:

- B1 radius (two-sided isotropic energy): $|\|z\| - \mathbb{E}\|z\||/\text{sd}$ in a PCA-64 space – the dilation/contraction energy axis.
- B2 RCF (Radial-Conditional, the novel core): a direction-conditioned radius $|\|z\| - \mathbb{E}[\|z\| \mid \text{dir}]|/\text{sd}$ via cosine- k NN; flags under-dispersion – the first branch to take contractions above chance per frame.
- B3 L4 Mahalanobis (one-sided) on the layer-4 $[\mu, \sigma^2, \kappa]$ block – the subspace axis for temporal_jitter.
- B4 spatial (when phi_spatial is extracted): Ledoit–Wolf Mahalanobis on the spatial-dispersion features.

Each branch is standardized on held-out clean and fused by a calibrated max. The shipped default folds an empirical-covariance Mahalanobis into B1 (radius') and replaces the plain max with a studentized combiner $\max - \alpha \cdot \text{median}$ ($\alpha \approx 0.5$); the derivation and ablation are in Appendices C–E. The detector is fit once on clean-train only and is corruption-blind; an optional per-recording aggregation integrates the consistent per-sequence bias.

4 Experimental Results

4.1 Setup

Prophesee Gen1 test split: 470 sequences, $\approx 343,099$ frames; each corruption applied per whole sequence at L1–L5. Detectors are fit on clean-train frames only and are corruption-blind; the split honours seq_lens so no clean frame appears in both fit and eval (vmem_utils.split_train_eval/held_out_eval). We report AUROC (and, in the appendix, AUPR / FPR@95%) per-frame and per-sequence. All AUROC values are oracle-labelled separability point estimates on a single split; bootstrap CIs are not yet banked (Section 5). Extraction is strictly $B = 1$ (SpikingJelly’s batch axis is the time axis).

4.2 Motivation: a SOTA detector collapses under event corruptions

We take the published, pretrained hybrid SNN–ANN detector (Neftci et al., 2025) (clean Gen1 test mAP ≈ 0.36) off the shelf and feed it corrupted input with no retraining, intercepting the event histogram after loading and before the backbone and applying each Gen1-C

Method	Required extra computation	Why infeasible on neuromorphic HW
Vmem- φ (ours)	native V register read + moments	— (it is free)
MSP (Hendrycks & Gimpel, 2017) / Energy (Liu et al., 2020)	softmax / log-sum-exp over logits	cores emit spikes, not float logits; no on-chip softmax
ODIN (Liang et al., 2018)	temperature softmax + input perturbation	perturbation needs an input-gradient pass the chip cannot take
ReAct (Sun et al., 2021)	activation clipping + second forward pass	re-forwarding violates the SNN energy budget
ViM (Wang et al., 2022)	SVD / dense linear algebra on features	no dense float linear-algebra unit; data movement too costly
GradNorm (Huang et al., 2021)	gradient backprop at inference	neuromorphic inference has no backward pass
kNN / Mahalanobis on ANN features (Lee et al., 2018; Sun et al., 2022)	an auxiliary ANN forward pass	requires a second, non-spiking network – defeats the SNN premise

Table 3: Standard OOD detectors each require a capability absent from the spiking, feed-forward, integer-event execution model; Vmem- φ requires none because the signal is the neuron state itself.

corruption at each severity. mAP is computed by the model’s native Prophesee evaluator, so the only change versus a stock validation run is the corruption of the input tensor. Reproducible from HybridDetection/validation_corrupt.py; Table 4 reports the result.

FIGURE/TABLE TODO: Insert results/neftci_map_degradation_table.tex once the GPU sweep is run. Rows = 6 corruptions, columns = severities 1–5, cells = mAP; clean baseline mAP ≈ 0.36 in the caption. (Modal/A100: needs the Gen1 test set + gen1_mAP36.ckpt.)

Table 4: Detection mAP of the pretrained hybrid detector under each corruption/severity (no retraining). Pending the GPU sweep. Hypotheses (from Section 3.3): hot_pixel/event_flood swamp object evidence; spatial_dropout deletes object regions; polarity_flip inverts the ON/OFF evidence the head was trained on (degrading mAP even though the membrane cannot flag it – a deliberate detectability-vs-harm mismatch).

Corruption	L1	L2	L3	L4	Corruption	L1	L2	L3	L4	L5
hot_pixel	0.983	1.000	1.000	1.000	hot_pixel	0.999	1.000	1.000	1.000	1.000
temporal_jitter	0.617	0.640	0.663	0.737	temporal_jitter	0.649	0.659	0.682	0.774	0.868
event_rate_shift	0.412	0.440	0.674	0.713	event_rate_shift	0.364	0.411	0.752	0.803	0.961
polarity_flip	0.415	0.422	0.437	0.459	polarity_flip	0.376	0.396	0.434	0.485	0.527
event_flood	0.413	0.418	0.428	0.445	event_flood	1.000	1.000	1.000	1.000	1.000
spatial_dropout	0.407	0.406	0.404	0.404	spatial_dropout	0.350	0.345	0.339	0.327	0.297

Table 5: Fused MDD per-frame AUROC. Table 6: Fused MDD per-sequence AUROC.

4.3 Main result: MDD OOD detection

Tables 5–6 report fused MDD AUROC per corruption/severity, per-frame and per-sequence (source: mdd_metrics.csv, mdd_metrics_aggregated.csv, generator evaluate_mdd.py). One unsupervised, corruption-blind score takes four of six corruptions to ≥ 0.96 with sequence aggregation; sequence aggregation rescues event_flood from 0.47 to 1.00 (the strongest frame→sequence jump) and lifts event_rate_shift to 0.96. The two residuals are theory-predicted dead ends.

What aggregation does (and does not) buy. Because each Gen1-C corruption is applied to the whole sequence, every frame of a corrupted recording carries the same bias; averaging the per-frame score over a recording integrates that consistent bias and separates the classes at a rate $\sim \sqrt{\#frames}$. This is why a near-chance per-frame signal (event_flood, 0.469) becomes perfectly separable per sequence (1.000): the lift is real but partly a property of the evaluation protocol, not evidence that any single frame is detectable. It would shrink if only a fraction of frames were corrupted, or if the decision had to be taken per frame. We therefore report per-frame AUROC as the conservative signal and treat per-sequence AUROC as an optimistic bound at a coarser (per-recording) decision granularity; the window sweep of Section 4.6 makes this dependence explicit, separating corruptions genuinely solved per frame from those solved only by pooling.

4.4 The geometry of corruption

The AUROC tables say how well each corruption is detected; Figure 3 shows why, making the dilation/contraction/invariant taxonomy visible. Membrane energy $\|z\|$ moves right for dilations and left for contractions (Figure 3a); the Mahalanobis d^2 of the contractions sits at smaller values than clean (Figure 3b), which is exactly the AUROC < 0.5 inversion; and the clean→corrupt centroid drift in a shared clean-PCA (Figure 3c) separates the corruption directions, with event_rate_shift and spatial_dropout sharing an axis (Appendix M, Figure 9).

4.5 Branch decomposition and fusion ablation

Figure 4a shows each branch is corruption-specialized (radius owns dilations, L4 owns timing, RCF/spatial own contractions); the plain calibrated max sits below the best single branch on both residuals, motivating the studentized fusion (Appendix D). Folding an empirical-




Figure 2: ROC curves of the improved fused MDD at L5 (per-frame), clean vs. each corruption. Source: plot_corruption_graphs.py (18_roc_curves_L5).




Figure 3: Geometry of corruption (L5, leakage-safe). (a) Dilations move energy right, contractions left, polarity overlaps clean. (b) Contractions sit at smaller d^2 than clean – the source of AUROC < 0.5. (c) Clean $\rightarrow$ corrupt centroid drift in a shared clean-PCA. Source: plot_corruption_graphs.py.

covariance Mahalanobis into the radius branch and replacing max with $\max - \alpha \cdot \text{median}$ (Figure 4b,c) recovers spatial_dropout from 0.560 toward its RCF ceiling (0.689) and lifts the worst-corruption floor, with no easy-corruption regression. The α knob trades polarity vs. dropout (Figure 4c).

4.6 Severity and sequence-aggregation sweep

Which results are manufactured by pooling? Figure 5 sweeps the aggregation window $W = 1 \dots \text{full}$. event_flood climbs $0.55 \rightarrow 1.00$ (crossing 0.9 at $W \approx 64$) and hot_pixel is saturated at every W , while polarity_flip ($0.54 \rightarrow 0.61$) and spatial_dropout ($0.54 \rightarrow 0.56$) stay flat – confirming they are the genuine residuals, not pooling failures. Read the other way, event_flood’s climb to 1.00 is largely manufactured by pooling: it has almost no per-frame signal and is “solved” only at the recording level, whereas hot_pixel is genuinely solved at every window. Pooling cannot rescue what is not present per frame – which is precisely why the flat residuals are the honest hard cases.




Figure 4: Branches and fusion (per-sequence, L5). (a) Per-branch AUROC heatmap. (b) Baseline max vs. improved (radius' + studentized fusion). (c) AUROC vs. studentization α : dropout $\uparrow$, polarity $\downarrow$. Source: plot_corruption_graphs.py. The combiner sweep is Figure 7.




Figure 5: Aggregation and severity. (a) Fused AUROC vs. pooling window per corruption. (b) Improved-MDD AUROC vs. severity at $W = 1 / 64 /$ full sequence. Source: plot_corruption_graphs.py.

4.7 Discussion: physical causes and biological adaptation

The value of the closed-form theory is that each physical axis (Section 3.2.1) maps to a distinct statistical axis of the membrane moments, so a detector's behaviour on a corruption is predictable from the sensor failure it models. One mechanism is worth dwelling on. `event_flood` – our worst per-frame corruption – inflates every channel's mean proportionally, mimicking a busy scene (Section 3.3). Biological neurons do not integrate a flood passively: spike-frequency adaptation (calcium- and slow- K^+ -driven) (Benda & Herz, 2003; Gabbiani & Krapp, 2014) and short-term synaptic depression (a divisive gain control) (Abbott et al., 1997; Rothman et al., 2009) normalise sustained drive so the neuron responds to changes, not absolute level. The PLIF neuron has none of this – a single static leak τ , a linear accumulator (Eq. 2), no adaptation – so a flood simply scales the accumulator and looks like a busy scene. The corruption biology handles best is the one our biologically-simplified model handles worst; equipping the front-end with adaptive (SFA / AdLIF) neurons is a concrete, biologically-grounded path to both detecting `event_flood` and restoring mAP (future work, cf. Section 4.2).

5 Limitations and the Open Challenge

Honest residuals. Two corruptions remain unsolved from the membrane alone, and the theory predicts why. `spatial_dropout`'s signature is spatial and GAP discards spatial layout; the released spatial-dispersion features (`phi_spatial`) feed MDD's spatial branch, which we validated end-to-end on a controlled contraction (a signal present only in `phi_spatial` is lifted from fused-without-spatial AUROC ≈ 0.50 to ≈ 1.00 ; Appendix L). `polarity_flip` is a property of the trained network's polarity invariance, not of our representation, and is undetectable from V_{mem} by construction. We frame both as an open challenge: the Achilles' heel of event-based perception, where corruptions are at once harmful (Section 4.2) and hard to detect cheaply.

Other limitations. (i) The cost claim is architectural (simulation, not silicon); on-chip energy is future work. (ii) AUROC are single-split point estimates without bootstrap CIs. (iii) Single dataset (Gen1); DSEC/Gen4 transfer is untested. (iv) Severities are synthetic monotone transforms, not measured sensor failures – the contribution is the benchmark/protocol/feature release. (v) Per-sequence AUROC assumes a corruption is consistent across a whole recording (as in Gen1-C); the strong per-sequence numbers (e.g. event_flood 1.000) are partly a consequence of that protocol and of the coarser decision granularity, and should not be read as per-frame detectability. (vi) AUROC is an oracle-labelled separability metric, not a value the corruption-blind detector reproduces at a fixed operating threshold in deployment.

6 Conclusion

We introduced OOD detection for event-camera spiking detectors from the free sub-threshold membrane potential, the Gen1-C corruption benchmark with its physical-mechanism grounding, a closed-form theory of per-corruption detectability, and the Manifold-Decomposition Detector that follows from a dilation/contraction diagnosis. A pretrained SOTA detector collapses under these corruptions; φ is the only OOD signal that runs inside the target hardware’s compute envelope; and one unsupervised score takes four of six corruptions to ≥ 0.96 with sequence aggregation. The two residuals are theory-predicted and released as an open challenge, with a concrete, biologically-grounded path – spatial-dispersion features and adaptive neurons – toward closing them.

References

- L. F. Abbott, J. A. Varela, Kamal Sen, and S. B. Nelson. Synaptic depression and cortical gain control. *Science*, 275(5297):221–224, 1997.
- Jan Benda and Andreas V. M. Herz. A universal model for spike-frequency adaptation. *Neural Computation*, 15(11):2523–2564, 2003.
- Mike Davies et al. Loihi: A neuromorphic manycore processor with on-chip learning. *IEEE Micro*, 38(1):82–99, 2018.
- Wei Fang, Zhaofei Yu, Yanqi Chen, Timothée Masquelier, Tiejun Huang, and Yonghong Tian. Incorporating learnable membrane time constant to enhance learning of spiking neural networks. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.
- Steve B. Furber, Francesco Galluppi, Steve Temple, and Luis A. Plana. The spinnaker project. *Proceedings of the IEEE*, 102(5):652–665, 2014.
- Fabrizio Gabbiani and Holger G. Krapp. Spike-frequency adaptation. *Scholarpedia*, 2014.
- Dan Hendrycks and Thomas Dietterich. Benchmarking neural network robustness to common corruptions and perturbations. In *International Conference on Learning Representations (ICLR)*, 2019.
- Dan Hendrycks and Kevin Gimpel. A baseline for detecting misclassified and out-of-distribution examples in neural networks. In *International Conference on Learning Representations (ICLR)*, 2017.
- Rui Huang, Andrew Geng, and Yixuan Li. On the importance of gradients for detecting distributional shifts in the wild. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- Kimin Lee, Kibok Lee, Honglak Lee, and Jinwoo Shin. A simple unified framework for detecting out-of-distribution samples and adversarial attacks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.

- Shiyu Liang, Yixuan Li, and R. Srikant. Enhancing the reliability of out-of-distribution image detection in neural networks. In International Conference on Learning Representations (ICLR), 2018.
- Weitang Liu, Xiaoyun Wang, John D. Owens, and Yixuan Li. Energy-based out-of-distribution detection. In Advances in Neural Information Processing Systems (NeurIPS), 2020.
- Aitor Martínez et al. A novel out-of-distribution detection approach for spiking neural networks: design, fusion, performance evaluation and explainability. Information Fusion, 2023.
- Emre Neftci et al. Efficient event-based object detection: A hybrid neural network with spatial and temporal processing. In IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2025.
- Jason S. Rothman, Laurence Cathala, Volker Steuber, and R. Angus Silver. Synaptic depression enables neuronal gain control. Nature, 457:1015–1018, 2009.
- Johannes Schemmel, Daniel Brüderle, Andreas Grünbl, Matthias Hock, Karlheinz Meier, and Sebastian Millner. A wafer-scale neuromorphic hardware system for large-scale neural modeling. In IEEE International Symposium on Circuits and Systems (ISCAS), pp. 1947–1950, 2010.
- Yiyu Sun, Chuan Guo, and Yixuan Li. React: Out-of-distribution detection with rectified activations. In Advances in Neural Information Processing Systems (NeurIPS), 2021.
- Yiyu Sun, Yifei Ming, Xiaojin Zhu, and Yixuan Li. Out-of-distribution detection with deep nearest neighbors. In International Conference on Machine Learning (ICML), 2022.
- Haoqi Wang, Zhizhong Li, Litong Feng, and Wayne Zhang. Vim: Out-of-distribution with virtual-logit matching. In IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2022.

Appendix

This appendix documents the complete experimental program – including the levers that did not work and the one result we retracted – so the headline claims can be audited against everything we tried. Every subsection names the generating script and source artifact.

A Implementation, Splits, and Reproducibility

Feature extraction. A VmemMonitor hooks all four MultiStepParametricLIFNode layers (64/128/256/256, 704 channels), captures $V \in \mathbb{R}^{T \times 1 \times C \times H \times W}$, GAPs over space, and reduces to $[\mu, \sigma^2, \kappa]$ per channel (φ , 2112-D). phi_spatial (1408-D) adds per-channel spatial variance and participation ratio, stored float32 (float16 overflows spatial_var under event_flood).

Constraint BATCH_SIZE = 1. SpikingJelly’s multi-step PLIF treats the batch axis as time; $B > 1$ leaks membrane state across sequences. Throughput scales by parallel workers (run_parallel_extract.py).

Protocol. Leakage-safe fit on clean-train only, corruption-blind; the split honours seq_lens. Metrics: AUROC, AUPR, FPR@95% (evaluate_detectors.py), per-frame and per-sequence. No bootstrap CIs yet. Config constants (benchmark_config.py): BATCH_SIZE=1, PHI_SAVE_EVERY=5, PLIF_LAYERS=None, TRAJ_SAVE_N=50.

Corruption	Static- φ Mahalanobis AUROC (L5)	Behaviour
hot_pixel	1.000	trivial
temporal_jitter	0.709	moderate
event_rate_shift	0.675	moderate; phase transition $\sim$ sev 3
polarity_flip	0.429	below chance – polarity-symmetric
event_flood	0.408	below chance – moments preserved
spatial_dropout	0.286	anti-detectable ($\rho = -1$)

Table 7: Reference detector (Mahalanobis on static φ , full data, L5). Source: performance_brief.md; generator evaluate_detectors.py.

Corruption	radius	rcf	l4	spatial	fused (max)
hot_pixel	1.000	0.989	1.000	1.000	1.000
event_rate_shift	0.900	0.834	0.506	0.540	0.952
temporal_jitter	0.513	0.393	0.933	0.724	0.947
event_flood	1.000	0.989	0.896	1.000	1.000
polarity_flip	0.486	0.552	0.573	0.639	0.609
spatial_dropout	0.391	0.693	0.403	0.411	0.560

Table 8: Per-branch and fused per-sequence AUROC at L5. Source: mdd_improvements.md §1; harness verified against evaluate_mdd_windows.py.

B Baseline OOD Detectors

We fit seven post-hoc detectors on clean φ (fit_detectors.py) and score them on every corrupted run (evaluate_detectors.py): Mahalanobis (Ledoit–Wolf; the reference detector), k -NN ($k = 5$), GMM (5), OCSVM (RBF), PCA (50), MLP-AE, RealNVP. Table 7 gives the reference behaviour that frames the benchmark: three corruptions sit below chance (informative but inverted – corrupted lies closer to the clean mean than held-out clean).

ANN baselines. A ResNet-18 on event-image / voxel-grid inputs (evaluate_ann_baselines.py); not on-chip-feasible, reported only as an upper reference.

FIGURE/TABLE TODO: Insert ANN-baseline AUROC table once the full grid is run.

C MDD: Per-Branch Decomposition

Table 8 reports each branch in isolation (per-sequence, L5). The plain calibrated max sits below the best single branch on both residuals (spatial 0.639 > 0.609; rcf 0.693 > 0.560) – it inflates the clean floor.

D Residual-Corruption Improvement Loop

Under a strict no-re-extraction, no-new-branch constraint we improved only the existing branches and the fusion rule (wired into analysis/mdd.py as the default, toggleable). (A) Radius covariance enrichment: $\text{radius}' = \max(\|z\| - \text{energy}, \text{emp_maha64})$, where emp_maha64 is an empirical (unshrunk) covariance Mahalanobis in PCA-64 that catches the covariance rotation a polarity flip induces (polarity 0.656 > spatial 0.639) without the dilation loss of full whitening (0.900 $\rightarrow$ 0.600 on event_rate_shift). (B) Studentized fusion: $\max - \alpha \cdot \text{median}$ ($\alpha \approx 0.5$) so a lone strong branch survives while diffuse clean noise cancels. Table 9 and Figure 6.

Honest ceiling. Neither residual reaches 0.8: spatial_dropout (contraction) caps ≈ 0.69 , polarity_flip (polarity-symmetric network) ≈ 0.66 . Breaking 0.8 needs information not in

Corruption	baseline (max)	radius' + (max − 0.5 med)	+(max − 1.25 med)
hot_pixel	1.000	1.000	1.000
event_rate_shift	0.952	0.956	0.948
temporal_jitter	0.947	0.952	0.966
event_flood	1.000	1.000	1.000
polarity_flip	0.609	0.607	0.560
spatial_dropout	0.560	0.613	0.689
mean	0.845	0.855	0.860
min (worst-corruption floor)	0.560	0.607	0.560

Table 9: Effect of the two modifications (per-sequence AUROC, L5). Source: mdd_improvements.md §3.




Figure 6: Improved − baseline per-sequence AUROC, per corruption × severity. Source: plot_corruption_graphs.py (22_improvement_by_severity).

the current φ (temporal $V(t)$, finer spatial pooling, the raw polarity channel) – out of scope under no-re-extraction.

Timeline. Table 10 records the exp1–exp10 trail, including dead ends.

E Fusion-Combiner Ablation

We swept combiners over the calibrated branches: max, mean, top-2 mean, p -norm, quantile-rank, studentized max − α {mean(rest), median, min}. Plain max sits below the best branch on residuals; rank/quantile helps residuals but collapses easy corruptions (event_flood 1.00 → 0.61); studentized max − α median is the Pareto winner (Figure 7). Supervised LR fusion buys nothing per-frame and fails to generalize (LR leave-one-out ≈ 0.49).

F Severity and Window Sweep (full)

Figure 8 gives the per-branch window sweep and the corruption × window heatmap (companion to Figure 5). The fused-vs-best-branch (oracle) gap is in paper_figures.md; the headline gap is event_flood’s per-frame → per-sequence jump (0.47 → 1.00).

G Representation Ablation

Is the sub-threshold membrane the right signal? representation_ablation.py scores membrane φ (full_membrane) and its $\mu/\sigma^2/\kappa$ and per-layer slices (L4 carries jitter); spike-rate

#	Finding
exp1	broad survey: flagged under-dispersion (dropout) + kurtosis (polarity) as the live signals
exp2	faithful harness: rcf 0.693 (dropout), spatial 0.639 (polarity) best singles; fused max dilutes
exp3	selective fusions: rcf+shrink-var $\rightarrow$ 0.704; spatial+kurt+L1 $\rightarrow$ 0.662; L1 retains polarity (0.630)
exp4	density inversion / sp_pr / rcf tuning: contraction confirmed but ≤ 0.56
exp5	broaden l4: polarity 0.573 $\rightarrow$ 0.634, flood-l4 0.896 $\rightarrow$ 1.000; fused flat $\Rightarrow$ fusion is the bottleneck
exp6	studentized max-rest: dropout 0.560 $\rightarrow$ 0.605, floor up
exp7	α -sweep max $-\alpha$ mean(rest): Pareto tension (dropout $\uparrow$, polarity $\downarrow$)
exp8	robust floor: max $-\alpha$ median best; $\alpha = 1.5 \Rightarrow$ dropout 0.693
exp9	empirical-cov PCA-64 Mahalanobis = 0.656 (best polarity branch)
exp10	fold emp_maha into radius + studentized-median: both improved, floor 0.560 $\rightarrow$ 0.607

Table 10: Residual-corruption experiment timeline. Source: mdd_improvements.md §4.




Figure 7: Fusion-combiner sweep across corruptions (per-sequence, L5). Source: plot_corruption_graphs.py (17_fusion_combiners).

/ spike-entropy; ANN feature (last_ann_gap) and logits (head_cls_L0_gap); and fused (concat + LogReg).

FIGURE/TABLE TODO: Insert the representation-comparison AUROC table from representation_ablation.py; PENDING the full-grid run.

H Severity Monotonicity, Cross-Corruption, and What φ Encodes

Spearman ρ . severity.py: hot_pixel $\rho = +1$, spatial_dropout $\rho = -1$ (anti-detectability).

FIGURE/TABLE TODO: Insert the per-corruption Spearman- ρ table.

Cross-corruption zero-shot. cross_corruption.py holds out one corruption type and evaluates on it.

FIGURE/TABLE TODO: Insert the cross-corruption transfer matrix.

Beyond binary OOD. analyse_comparisons.py trains, on φ , a 7-class corruption-type classifier (confusion matrix) and a severity regression (R^2).

FIGURE/TABLE TODO: Insert the confusion matrix and severity-regression R^2 table.




Figure 8: Window sweep detail. Source: plot_corruption_graphs.py.

Lever	Result
Two-sided / folded Mahalanobis	rescues the 3 below-chance (dropout 0.286 $\rightarrow$ 0.531, flood 0.408 $\rightarrow$ 0.577, polarity 0.429 $\rightarrow$ 0.547) but degrades jitter (0.709 $\rightarrow$ 0.615)
Global activity scalar (two-sided σ^2 energy)	big win for event_rate_shift (0.675 $\rightarrow$ 0.850); near-useless elsewhere
More fit samples (3k $\rightarrow$ 50k)	free but small: jitter 0.709 $\rightarrow$ 0.757; rest flat
Per-feature standardization before Ledoit–Wolf	jitter 0.709 $\rightarrow$ 0.820 but hurts event_rate_shift (0.675 $\rightarrow$ 0.568)
Naive max fusion (two-sided + activity)	worse than either alone
Temporal AE (1D-CNN)	weak everywhere (0.40–0.58)
Handcrafted temporal (leakage-safe)	modest win on flood/dropout ($\approx$ 0.54–0.62)
Supervised LR meta-fusion	no per-frame gain; LOO $\approx$ 0.49

Table 11: Levers tried on the hard corruptions. Oracle best-achievable per corruption: hot_pixel 1.00, event_rate_shift 0.85, temporal_jitter 0.82–0.84, event_flood 0.63 (frame), polarity_flip 0.55, spatial_dropout 0.54 (frame). Source: performance_brief.md §6.

I Temporal and Sequence Representations (and a Retracted Result)

We evaluated handcrafted temporal_phi (28-D), a Temporal-AE (1D-CNN on the GAP’d $V(t)$ trajectory), and margin histograms ($V - \theta$ quantized).

Retraction (honest negative). An earlier internal finding claimed temporal features rescue the hard corruptions to ≈ 0.85 and that “the paper’s strongest version is a temporal paper.” This does not reproduce. It was computed on the 50-sample raw-trajectory path (35 train/15 test from the first 50 frames of one sequence, a 28-D Mahalanobis on 35 autocorrelated points, without a held-out clean split). Re-running the exact method on freshly extracted data gives 0.29/0.04/0.06 – the 0.85 was leakage plus 50-sample noise (a refactor: unified train/eval split commit postdates the finding).

Honest picture. With a leakage-safe split, handcrafted temporal beats static only modestly: event_flood ≈ 0.625 (static ≈ 0.58), spatial_dropout ≈ 0.541 (static ≈ 0.47) – a real ~ 0.05 –0.08 gain, not a rescue. The Temporal-AE is the weakest detector. The genuine contraction lift comes from sequence aggregation (Section 4.6).

J Levers Tried (Negative-Results Catalogue)

K Free-Rider Validity Check

free_rider_ablation.py compares OOD AUROC from (A) trained-SNN membrane, (B) random-init SNN membrane, (C) raw-input statistics; the premise requires (A) $\gg$ (C).




Figure 9: Additional diagnostics. (a) cosine similarity of mean-shift directions (rate_shift ↔ dropout cluster); (b) spatial_var/spatial_pr clean vs. flood vs. dropout; (c) variance-block shift (dropout sits low = under-dispersion); (d) per layer/moment standardized channel shift; (e) shared clean-PCA scatter; (f) log d^2 violins (contraction medians dip toward clean). Source: plot_corruption_graphs.py.

FIGURE/TABLE TODO: Insert the trained-vs-random-vs-raw table; PENDING run – a gating experiment.

L Spatial-Branch Mechanism Check

On a controlled contraction whose signal lives only in phi_spatial (GAP'd φ matched to clean): fused-without-spatial AUROC = 0.496 → fused-with-spatial = 1.000; spatial branch alone = 1.000 – confirming the wiring. On the banked extraction the spatial branch nails event_flood (1.000) but stays below chance on spatial_dropout (0.411), consistent with the contraction ceiling.

M Additional Geometry Figures

Figure 9 collects the remaining geometry/diagnostic panels.

N Released Artifacts: Gen1-C

Corruption benchmark (code + seeds): six corruptions × five severities, fully seeded; a (corruption, severity, seed) triple regenerates the exact stream via apply_corruption(..., rng); NumPy/CuPy-agnostic on the standard $(N, 20, H, W)$ histogram. We ship the library + seed manifest + config that regenerate from the user's own Gen1 copy (not redistributed histograms). Released φ feature set: pre-extracted membrane statistics for clean + 30 corrupted runs (with seq_lens) so detectors can be benchmarked without GPU inference. Protocol: the leakage-safe split, per-frame / per-sequence AUROC, and the 7-detector + MDD harness. Action before release: confirm the Prophesee Gen1 license permits distributing derived features / regeneration code.